



Basic Search Algorithms

CSE 415: Introduction to Artificial Intelligence
University of Washington
Winter, 2019

© S. Tanimoto and University of Washington, 2019



Outline

- Recursive Depth-First Search
- Graph Search
- Iterative Depth-First Search
- Breadth-First Search
- Iterative Deepening
- Graphs with Edge Costs
- Uniform-Cost Search
- Heuristics and Best-First Search
- A* Search

CSE 415, Univ. of Wash.


Basic Search Algorithms

2



Painted-Squares Puzzle

In the previous lecture, we introduced a simple puzzle to illustrate tree search. Now let's formulate the search algorithm.



Tree of states for a
2x2 Painted
Squares puzzle

CSE 415, Univ. of Wash.

State Spaces

3



Recursive Depth-First Method*

Current board $B \leftarrow$ empty board.
Remaining pieces $Q \leftarrow$ all pieces.
Call $\text{Solve}(B, Q)$.

Procedure $\text{Solve}(\text{board } B, \text{set of pieces } Q)$

For each piece P in Q , {

For each orientation A {

Place P in the first available position of B in orientation A , obtaining B' .

If B' is full and meets all constraints, output B' .

If B' is full and does *not* meet all constraints, return.

Call $\text{Solve}(B', Q - \{P\})$.

}

}

*Sometimes covered in CSE 143 or CSE 373. Also known as Recursive Backtracking Search.

CSE 415, Univ. of Wash.

Basic Search Algorithms

4



Graph Search

exploring none-tree graphs. Careful about cycles causing infinite loops! MUST keep track of what we've done so far (i.e. in DFS which vertices we visited)

When descendant nodes can be reached with moves via two or more paths, we are really searching a more general graph than a tree.

Whether we are in a tree or a more general graph, the basic idea of depth-first search is the same... do not backtrack until all descendants have been tried.

Depth-First Search: Examine the nodes of the graph by fully exploring the "descendants" of a node before trying any "siblings" of a node.


CSE 415, Univ. of Wash.

Basic Search Algorithms

5



Sample Graph



CSE 415, Univ. of Wash.

Basic Search Algorithms

6



Depth-First Search: Iterative Formulation

1. Put the start state on a list OPEN
2. If OPEN is empty, output "DONE" and stop.
3. Select the first state on OPEN and call it S.
Delete S from OPEN.
Put S on CLOSED.
If S is a goal state, output its description
4. Generate the list L of successors of S and delete from L those states already appearing on CLOSED.
5. Delete from OPEN any members of OPEN that occur on L.
Insert all members of L at the front of OPEN.
6. Go to Step 2.

OPEN = states/vertices reached but not fully explored
CLOSED = vertices already explored
i.e. open = visited from a neighbor and added to the list of next vertex
closed = already explored all neighbors and marked them to visit.

CSE 415, Univ. of Wash.

Basic Search Algorithms

7



Breadth-First Search: Iterative Formulation

1. Put the start state on a list OPEN
2. If OPEN is empty, output "DONE" and stop.
3. Select the first state on OPEN and call it S.
Delete S from OPEN.
Put S on CLOSED.
If S is a goal state, output its description
4. Generate the list L of successors of S and delete from L those states already appearing on CLOSED.
5. Delete from L any members of OPEN that occur on L.
Insert all members of L at the end of OPEN.
using OPEN list as Stack = DFS; Queue = BFS
6. Go to Step 2.

possibly path to that vertex as well → have nodes store info of which neighbor they were added to OPEN by.

CSE 415, Univ. of Wash.

Basic Search Algorithms

8



Graph Search vs. Tree Search

Note that these 2 formulations are for searching general graphs.

If we know we are doing tree search, then we don't need to worry about deleting anything in steps 4 and 5, and we don't need a CLOSED list.

CSE 415, Univ. of Wash.

Basic Search Algorithms

9



Comparing DFS and BreadthFS

Optimality of the solution path: this is because BFS moves one layer at a time.

When BreadthFS arrives at a goal node, a shortest (i.e., optimal) path is easily extracted by "backtracing" (not backtracking).

However, DFS may arrive at a goal node via a very non-optimal path, and the shortest path is not readily available at that point.

Memory Utilization:

BreadthFS usually requires a large amount of memory. For example, in a tree with branching factor b , the OPEN list will require storage $O(b^d)$, where d is the depth of the closest goal node.

On the other hand, DFS requires only $O(bd)$ storage for its OPEN list.

(No CLOSED lists are needed if we know that the problem-space graph is acyclic.)

CSE 415, Univ. of Wash.

Basic Search Algorithms

10



compromising DFS

Iterative Deepening DFS

memory properties of DFS (does not hog as much space) and gives optimal path like BFS. However, recomputes operations, so slower than BFS

We can combine the benefits of DFS and BreadthFS to get optimal paths without huge memory requirements.

Instead of regular BreadthFS, we do a sequence of DFS executions, but with a depth limit for each execution. We make the depth limit increase by 1 in each execution, starting from 0.

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or failure
  inputs: problem, a problem
  for depth ← 0 to ∞ do
    result ← DEPTH-LIMITED-SEARCH(problem, depth)
    if result ≠ cutoff then return result
```

Russell & Norvig

keep running DFS and increasing the maximum depth allowed.

CSE 415, Univ. of Wash.

Basic Search Algorithms

11



Is Iterative Deepening DFS Efficient?

IDDFS saves memory, but at the expense of extra time, due to repeated searching of the upper tree levels.

Assuming we are searching a tree with an average branching factor $b > 1$, then an analysis of the number of repeat visits to nodes leads us to conclude "Yes, it's efficient."

The repeated work is bounded by a small factor (effectively constant, because it's independent of the size of Σ).

Russell & Norvig

CSE 415, Univ. of Wash.

Basic Search Algorithms

12

Iterative Deepening DFS Example
 $\angle = 0$

Limit = 0

Russell & Norvig

CSE 415, Univ. of Wash. Basic Search Algorithms 13

IDDFS Example (cont.)
 $\angle = 1$

Limit = 1

Russell & Norvig

CSE 415, Univ. of Wash. Basic Search Algorithms 14

IDDFS Example (cont.)
 $\angle = 2$

Limit = 2

Russell & Norvig

CSE 415, Univ. of Wash. Basic Search Algorithms 15

IDDFS Example (cont.)
 $\angle = 3$

Limit = 3

Russell & Norvig

CSE 415, Univ. of Wash. Basic Search Algorithms 16

Overhead for Iterative Deepening

Repeated work takes place mainly near the root, where there are relatively few nodes.

With $b = 2$, the overhead is less than a factor of 2. (e.g., 57/31)

Depth	N in level	N in tree	IDDFS
0	1	1	1
1	2	3	4
2	4	7	11
3	8	15	26
4	16	31	57

$2^{d+1} - 1$
 $\text{sum}(2^{i+1} - 1) \text{ } i = 0 \text{ to } d$
 $= 2(2^{d+1} - 1) - (d+1)$
 $= 2^{d+2} - (d+3)$

IDDFS is the cumulative sum of N in tree

Russell & Norvig

CSE 415, Univ. of Wash. Basic Search Algorithms 17

Search Algorithms

Alternative objectives:

- Reach any goal state
- Find a short or shortest path to a goal state

Alternative properties of the state space and moves:

- Tree structured vs graph structured, cyclic/acyclic
- Weighted/unweighted edges

Alternative programming paradigms:

- Recursive
- Iterative
- Iterative deepening
- Genetic algorithms

Russell & Norvig

CSE 415, Univ. of Wash. Basic Search Algorithms 18



Problem-Space Graphs with Weighted Edges

Let Σ be space of possible states, and S be the corresponding set of nodes. Let (s_i, s_j) be an edge representing a move from s_i to s_j . $w(s_i, s_j)$ is the weight or **cost** associated with moving from s_i to s_j .

The **cost of a path** $[(s_1, s_2), (s_2, s_3), \dots, (s_{n-1}, s_n)]$ is the sum of the weights of its edges.

A **minimum-cost path** P from s_1 to s_n has the property that for any other path P' from s_1 to s_n , $\text{cost}(P) \leq \text{cost}(P')$.


CSE 415, Univ. of Wash.

Basic Search Algorithms

19



Graphs with Weighted Edges



CSE 415, Univ. of Wash.

Basic Search Algorithms

20



Uniform-Cost Search

A more general version of breadth-first search.

Processes states in order of increasing path cost from the start node.

The list OPEN is maintained as a priority queue. Associated with each node is its current best estimate of its distance from the start state.

As a node s_i from OPEN is processed, its successors are generated. The tentative distance for a successor s_j of node s_i is computed by adding $w(s_i, s_j)$ to the distance for s_i .

If s_j occurs on OPEN, the smaller of its old and new distances is retained. If s_j occurs on CLOSED, and its new distance is smaller than its old distance, then it is taken off of CLOSED, put back on OPEN, but with the new, smaller distance.

CSE 415, Univ. of Wash.

Basic Search Algorithms

21



Heuristics

A heuristic is a "rule of thumb" for operating in unknown, uncertain, or complex environments or problem-solving contexts.

A heuristic evaluation function, in state-space search, is a function $h: \Sigma \rightarrow \mathbb{R}^+$ that can be used as an estimate of how close a state is to a goal or simply to prioritize states for attention.

Examples:

Euclidean distance between a city and the goal. (in the routing problem)

Number of pieces not yet placed in a puzzle. (painted squares).

Average distance a puzzle piece (in the 8-puzzle) has to move on the board to get to its destination.

Hot-cold (in a game of Find-the-hidden-object). Hot: close to 0.

Cold: much greater than 0.

CSE 415, Univ. of Wash.

Basic Search Algorithms

22



Best-First Search

Provided we have a heuristic evaluation function, we can prioritize states for expansion using the function.

By changing our iterative formulation of **Depth-First Search** to use a **PRIORITY QUEUE** to implement the OPEN list, we get **Best-First Search**.

CSE 415, Univ. of Wash.

Basic Search Algorithms

23



Ideal Distances in A* Search

Let $f(s)$ represent the cost (distance) of a shortest path that starts at the start node, goes through s , and ends at a goal node.

Let $g(s)$ represent the cost of a shortest path from the start node to s .

Let $h(s)$ represent the cost of a shortest path from s to a goal node.

Then $f(s) = g(s) + h(s)$

During the search, the algorithm generally does not know the true values of these functions.

CSE 415, Univ. of Wash.

Basic Search Algorithms

24



Estimated Distances in A* Search

Let $g'(s)$ be an estimate of $g(s)$ based on the *currently known shortest distance* from the start node to s .

Let the $h'(s)$ be a *heuristic* evaluation function that estimates the distance (path length) from s to the nearest goal node.

Let $f'(s) = g'(s) + h'(s)$

Best-first search using $f'(s)$ as the evaluation function is called *A* search*.

CSE 415, Univ. of Wash.

Basic Search Algorithms

25



Admissibility of A* Search

Under certain conditions, A* search will always reach a goal node and be able to identify a shortest path to that node as soon as it arrives there.

The conditions are:

$h'(s)$ must not exceed $h(s)$ for any s .

$w(s_i, s_j) > 0$ for all s_i and s_j .

This property of being able to find a shortest path to a goal node is often called the admissibility property of A* search. However, we should consider admissibility as a property of h' . The heuristic h' is admissible if and only if it satisfies $0 \leq h'(s) \leq h(s)$, for all nodes s .

Sometimes we say that a particular A* algorithm is admissible. We can say this when its h' function satisfies the admissibility condition and the underlying search problem involves only positive weights.

CSE 415, Univ. of Wash.

Basic Search Algorithms

26



Search Algorithm Summary

	Unweighted graphs	Weighted graphs
blind search	Depth-first Breadth-first IDDFS	Depth-first Uniform-cost Iterative Lengthening
uses heuristics	Best-first	A* Iterative Deepening A*

CSE 415, Univ. of Wash.

Basic Search Algorithms

27