

Ss
State-space
Search

START LECT WED JAN 23

Heuristic Search

CSE 415: Introduction to Artificial Intelligence
University of Washington
Winter, 2019

© S. Tanimoto and University of Washington, 2019

Ss
State-space
Search

Outline

Motivation, Definition
A* Algorithm
Admissibility and Consistency
Heuristics for the 8 Puzzle
Designing Heuristics

CSE 415, Univ. of Wash. Heuristic Search

2

Ss
State-space
Search

Motivation

DFS, BFS, IDDFS; do not consider qualities of states on their search
only keeps checking if they are at a goal state

- Blind search can waste time and space.
(due to the combinatorial explosion.)
- Additional knowledge MAY be available -
- how could it help?
 - E.g., finding a shortest route from Wash. U. to U. Wash. (St. Louis to Seattle).
 - Blind search considers all directions equally, including towards Wash. D.C.
 - Additional knowledge: we need to head northwest.

using additional knowledge to make INFORMED decisions about which states to transition to next; BIAS the search

CSE 415, Univ. of Wash.

Heuristic Search

3

Ss
State-space
Search

Starting from St. Louis--Blind



CSE 415, Univ. of Wash.

Heuristic Search

4

Ss
State-space
Search

BIASED SEARCH Starting from St. Louis--Informed



CSE 415, Univ. of Wash.

Heuristic Search

5

Ss
State-space
Search

Definition

reminder: Sigma = State space

A *heuristic function* (or simply heuristic) is a function $h: \Sigma \rightarrow \mathbb{R}$, that takes a state as its argument and returns a real number that is an estimate of the distance (or cost) from that state to the closest (or having lowest-cost path) goal state.

$$h(s) = r$$

The function h is typically based on *partial information* about the relationship between each state s and the closest goal state γ to s .

For example, if each state has an (x,y) location, then knowing only x_s and x_γ , we could estimate the distance between s and γ as $|x_s - x_\gamma|$.

typically, $h(\text{goal state}) = 0$

CSE 415, Univ. of Wash.

Heuristic Search

6



Outline

Motivation, Definition

A* Algorithm

Admissibility and Consistency

Heuristics for the 8 Puzzle

Designing Heuristics

CSE 415, Univ. of Wash.

Heuristic Search

7



A* Algorithm

variation on Dijkstra's algorithm that incorporates heuristic?

Given a state space Σ having a distance (or cost) function on moves (graph edges): $d(s_i, s_j)$, the A* algorithm searches for a shortest (lowest-cost) path from the initial state s_0 to a goal state γ .

The following algorithm gives the general control structure for A*. It omits a few details:

1. Back pointers for backtracing a path when a goal state is reached.
2. Details of computing g . (done in a manner similar to that in Dijkstra's algorithm, i.e., Uniform Cost Search).
3. Details of implementing the OPEN list and its methods for inserting, finding, and removing.

CSE 415, Univ. of Wash.

Heuristic Search

8



A* Algorithm

1. For the start state s_0 , compute $f(s_0) = g(s_0) + h(s_0) = h(s_0)$ and put $[s_0, f(s_0)]$ on a list OPEN.
2. If OPEN is empty, output "DONE" and stop.
3. Find and remove the item $[s, p]$ on OPEN having **highest priority** (lowest p). Break ties arbitrarily. Put $[s, p]$ on CLOSED.
4. If s is a goal state: output its description (and backtrace a path), and if h is known to be admissible, halt.
4. Generate the list L of $[s', f(s')]$ pairs where the s' are the successors of s and their f values are computed using $f(s') = g(s') + h(s')$. Consider each $[s', f(s')]$:
 - If there is already a pair $[s', q]$ on CLOSED (for any value q):
 - If $f(s') > q$, then remove $[s', f(s')]$ from L .
 - If $f(s') \leq q$, then remove $[s', q]$ from CLOSED. **we found a shorter path to s' ; must recompute all successors**
 - Else if there is already a pair $[s', q]$ on OPEN (for any value q):
 - If $f(s') > q$, then remove $[s', f(s')]$ from L .
 - If $f(s') \leq q$, then remove $[s', q]$ from OPEN.
5. Insert all members of L onto OPEN.
6. Go to Step 2.

CSE 415, Univ. of Wash.

Heuristic Search

9



A* Algorithm Behavior

During the search A* gives highest priority to that as-yet unexplored state (except in cases where some previously explored state needs to be re-examined) that has the lowest sum of distance from the initial state plus estimated distance to a goal.


CSE 415, Univ. of Wash.

Heuristic Search

10



Example



state s : s_0 A B C D E F G H γ
 heuristic $h(s)$: 14 15 4 10 3 2 16 10 5 0

We show newly enqueued $[s, p]$ pairs in green, and updated $[s, p]$ pairs in red.


CSE 415, Univ. of Wash.

Heuristic Search

11



Example




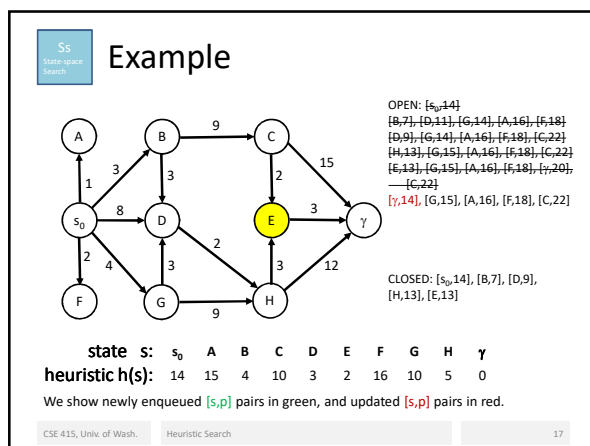
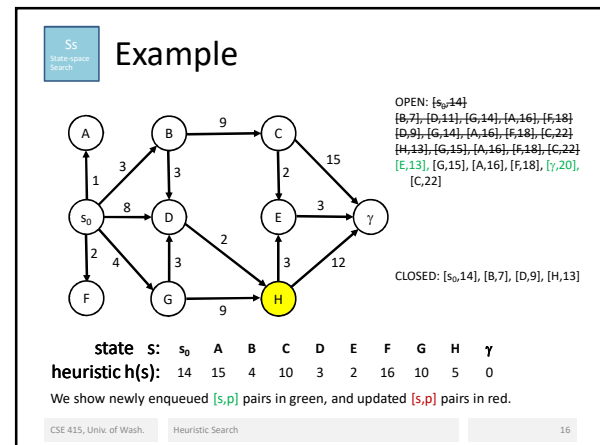
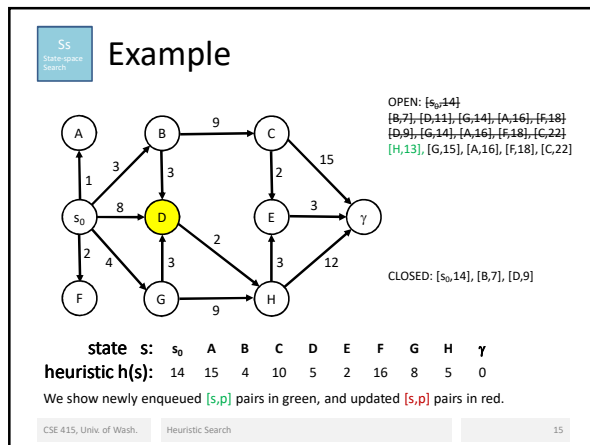
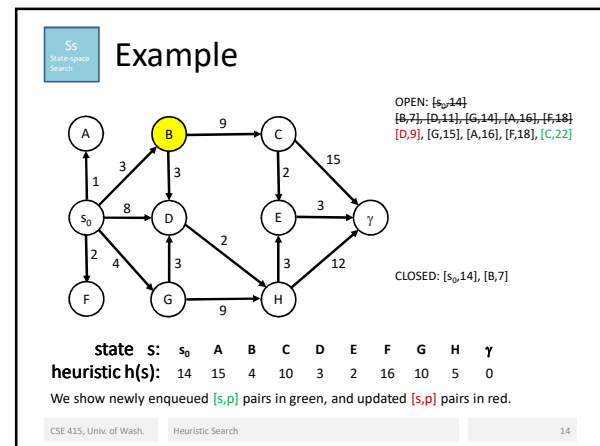

state s : s_0 A B C D E F G H γ
 heuristic $h(s)$: 14 15 4 10 3 2 16 10 5 0

We show newly enqueued $[s, p]$ pairs in green, and updated $[s, p]$ pairs in red.

CSE 415, Univ. of Wash.

Heuristic Search

12





Notes on the Example

The path found is the same as that found by UCS (Dijkstra).
However, fewer nodes are expanded; e.g., A and F are never expanded.

The heuristic h is admissible, so as soon as γ becomes the current state, we can stop.

But h is not consistent: $h(G) - h(D) > d(G, D)$.
This fact doesn't cause trouble here, fortunately.
(If h is consistent, no state will ever have to be expanded a second time, i.e., never have to be moved from CLOSED back to OPEN.)

CSE 415, Univ. of Wash.

Heuristic Search

19



Outline

Motivation, Definition

A* Algorithm

Admissibility and Consistency

Heuristics for the 8 Puzzle

Designing Heuristics

CSE 415, Univ. of Wash.

Heuristic Search

20

Note heuristic MUST be carefully chosen to ensure our path is actually minimal! The properties our heuristic must satisfy our Admissibility and Consistency



Admissibility and Consistency

Heuristic h is admissible if and only if
 $(\forall s \in \Sigma)(\forall \gamma \in \Gamma) h(s) \leq d(s, \gamma)$

Here $d(s, \gamma)$ is the length (cost) of the shortest (lowest-cost) path from s to γ .

I.e., h estimates but never overestimates the distance from s to the closest goal.

If h is admissible, then A*, using h , finds a shortest path (optimal path) as soon as it expands any goal γ .

Admissibility. As soon as we first reach a goal state, we have reached it via a minimal path

CSE 415, Univ. of Wash.

Heuristic Search

21



Admissibility and Consistency

Heuristic h is consistent if and only if
For each edge (s_i, s_j) in the problem-space graph,
 $h(s_i) - h(s_j) \leq d(s_i, s_j)$

Here $d(s_i, s_j)$ is the length (cost) of the edge from s_i to s_j .

If h is consistent, then along any shortest path from a node (state) s to its closest goal, then h values will be monotonically non-increasing along the path.

If h is consistent, then A*, using h , never has to re-expand a node.
i.e. remove a node from CLOSED and re-expand it

CSE 415, Univ. of Wash.

Heuristic Search

22

Consistency: never have to visit and expand a state more than once
i.e. ensures our search is efficient



Admissibility of A*

A* is *admissible*: Provided its heuristic is admissible, and a path exists from s_0 to γ , then A* will find a shortest path and stop.

CSE 415, Univ. of Wash.

Heuristic Search

23



Outline

Motivation, Definition

A* Algorithm

Admissibility and Consistency

Heuristics for the 8 Puzzle

Designing Heuristics

CSE 415, Univ. of Wash.

Heuristic Search

24

END LECT WED JAN 23

Heuristics for the Eight Puzzle

$h_0(s) = 0$ (uninformed; blind search)

$h_1(s)$ = number of tiles out of place. ("Hamming") $0 \leq h_1(s) \leq 8$.

$\gamma =$

	1	2
3	4	5
6	7	8

 goal state

$s_a =$

3	1	2
7	6	5
4		8

$h_1(s_a) = 4$.

h_1 is admissible IF we ignore the blank space, otherwise its not (consider the final move; both blank space and one tile would be out of place)

h_1 easy to compute

CSE 415, Univ. of Wash. Heuristic Search 25

Heuristics for the Eight Puzzle

$h_2(s)$ = number of rows tile 7 is away from its place.

$\gamma =$

	1	2
3	4	5
6	7	8

$s_a =$

3	1	2
7	6	5
4		8

$h_2(s_a) = 1$.

Using tile 7 is arbitrary here -- just an example of a heuristic based on a single, particular tile.

CSE 415, Univ. of Wash. Heuristic Search 26

Heuristics for the Eight Puzzle

$h_3(s)$ = number of rows **and** columns tile 7 is away from its place.

$\gamma =$

	1	2
3	4	5
6	7	8

$s_a =$

3	1	2
7	6	5
4		8

$h_3(s_a) = 2$.

This sum is known as the *Manhattan distance* (for a single tile).

Admissible AND dominates h_2 ; more informed = better BUT harder to compute

CSE 415, Univ. of Wash. Heuristic Search 27

Heuristics for the Eight Puzzle

$h_4(s)$ = sum of Manhattan distances for all 8 tiles.

$\gamma =$

	1	2
3	4	5
6	7	8

$s_a =$

3	1	2
7	6	5
4		8

$h_4(s_a) = 7$.

This is called the Manhattan distance heuristic.

In this example $h_4(s_a) = h(s_a)$ (the actual shortest distance).

Admissible, but again hard to compute; we see this trade off between the ease to compute and how informed the heuristic is quite often

CSE 415, Univ. of Wash. Heuristic Search 28

Heuristics for the Eight Puzzle

$h_5(s)$ = sum of Euclidean distances for all 8 tiles.

$$= \sum_{i=1}^8 \sqrt{dx_i^2 + dy_i^2}$$

$\gamma =$

	1	2
3	4	5
6	7	8

$s_a =$

3	1	2
7	6	5
4		8

$h_5(s_a) = 1 + 3\sqrt{2} \approx 5.2326$

This is called the Euclidean distance heuristic.

In (at least) this example $h_5(s_a) < h_4(s_a)$.

Euclidean is not as good as Manhattan.

CSE 415, Univ. of Wash. Heuristic Search 29

Heuristic Domination

If $(\forall s \in \Sigma) h_i(s) \geq h_j(s)$, then we say h_i dominates $h_j(s)$. However, we assume both heuristics are admissible.

If h_i dominates h_j , then we call h_i "more informed" than h_j . Having a highly informed heuristic is good for limiting a search to relevant parts of the state space.

However, one has to trade off this off against the higher computational cost that usually goes with more informed heuristics.

CSE 415, Univ. of Wash. Heuristic Search 30



Outline

Motivation, Definition
A* Algorithm
Admissibility and Consistency
Heuristics for the 8 Puzzle
Designing Heuristics

CSE 415, Univ. of Wash.

Heuristic Search

31



Designing Heuristics

A common approach to defining heuristics is to create a simpler type of problem. *remove some of the constraints...*

E.g., For the 8 Puzzle, allow tiles to be removed and put back anywhere, and "charge" a cost of 0.5 for each removal and 0.5 for each putting back. That leads to the Hamming heuristic.

Or: allow tiles to be piled up on top of one another, thus making it easier to move each tile (still one square at a time) to its destination. This leads to the Manhattan heuristic.

CSE 415, Univ. of Wash.

Heuristic Search

32



Designing Heuristics (cont)

Another way to simplify: change some of the tiles into "blanks" (like the blank tile in Scrabble). The new goal is to get only the non-blank tiles into their proper positions; the blanks are "don't-care" tiles that still take up space, but whose relative ordering is not important. For example:

$$\gamma' = \begin{bmatrix} & 1 & 2 \\ 3 & 4 & 5 \\ / & / & / \end{bmatrix}$$

$$s_a' = \begin{bmatrix} 3 & 1 & 2 \\ / & / & 5 \\ 4 & & / \end{bmatrix}$$

To compute $h(s_a')$, we transform s_a' into s_a'' , and solve the simplified problem, getting a path length $d(s_a'', \gamma')$, which we use as the value of h . If the reduced problem is easy enough, then we can precompute a table of $d(s_a'', \gamma')$ values to speed up computing h during the search. Such a table is called a *pattern database*.

CSE 415, Univ. of Wash.

Heuristic Search

33



Outline

Motivation, Definition
A* Algorithm
Admissibility and Consistency
Heuristics for the 8 Puzzle
Completeness and Optimality
Designing Heuristics
done!

CSE 415, Univ. of Wash.

Heuristic Search

34