



Uncertainty in AI: The Joint Distribution

CSE 415: Introduction to Artificial Intelligence
University of Washington
Winter, 2019

© S. Tanimoto and University of Washington, 2019



Outline

- The Monty Hall Problem revisited
- Joint probability distributions
- Marginal distributions
- Factored joint probability distributions
- Bayes nets
- Benefits of Bayes nets for expert systems


Univ. of Wash.

The Joint Distribution

2



The Monty Hall Problem (revisited)



Univ. of Wash.

The Joint Distribution

3



Joint Probability Distribution for the Monty Hall Problem

| Prize in | You choose | Host opens | P | Payoff if no switch | Payoff if switch |
|----------|------------|------------|------|---------------------|------------------|
| R | R | G | 1/18 | 1 | 0 |
| R | R | B | 1/18 | 1 | 0 |
| R | G | B | 1/9 | 0 | 1 |
| R | B | G | 1/9 | 0 | 1 |
| G | R | B | 1/9 | 0 | 1 |
| G | G | R | 1/18 | 1 | 0 |
| G | G | B | 1/18 | 1 | 0 |
| G | B | R | 1/9 | 0 | 1 |
| B | R | G | 1/9 | 0 | 1 |
| B | G | R | 1/9 | 0 | 1 |
| B | B | G | 1/18 | 1 | 0 |
| B | B | R | 1/18 | 1 | 0 |

Univ. of Wash.

The Joint Distribution

4



Discussion

Marginal probability of winning, never switching: $1/3$
Marginal probability of winning, always switching: $2/3$

Other marginal probabilities:

$P(\text{prize is behind Red door}) = 1/3$
 $P(\text{you choose Red door}) = 1/3$, assuming you choose randomly.
 $P(\text{you first choose the right door}) = 1/3$

The joint probability distribution gives us the means to answer many questions about random variables and their relationships.

Univ. of Wash.

The Joint Distribution

5



Another Joint Distribution

| Solar storm | Mike's battery on the fritz | Jan's radio works just fine | Mike's radio reception | Joint prob. |
|-------------|-----------------------------|-----------------------------|------------------------|-------------|
| F | F | F | bad | 0.0156 |
| F | F | F | good | 0.0468 |
| F | F | F | none | 0.0156 |
| F | F | T | bad | 0.1404 |
| F | F | T | good | 0.4212 |
| F | F | T | none | 0.1404 |
| F | T | F | bad | 0.0006 |
| F | T | F | good | 0.0002 |
| F | T | F | none | 0.0012 |
| F | T | T | bad | 0.0054 |
| F | T | T | good | 0.0018 |
| F | T | T | none | 0.0108 |

Univ. of Wash.

The Joint Distribution

6

Jd
joint
Distribution

(continued)

| Solar storm | Mike's battery on the fritz | Jan's radio works just fine | Mike's radio reception | Joint prob. |
|-------------|-----------------------------|-----------------------------|------------------------|-------------|
| T | F | F | bad | 0.0585 |
| T | F | F | good | 0.0117 |
| T | F | F | none | 0.0468 |
| T | F | T | bad | 0.039 |
| T | F | T | good | 0.0078 |
| T | F | T | none | 0.0312 |
| T | T | F | bad | 0.0006 |
| T | T | F | good | 0.00015 |
| T | T | F | none | 0.00225 |
| T | T | T | bad | 0.0004 |
| T | T | T | good | 0.0001 |
| T | T | T | none | 0.0015 |

Univ. of Wash.

The Joint Distribution

7

Jd
joint
Distribution

Bayes Nets

A practical way to manage probabilistic inference when multiple variables (perhaps many) are involved.

Requirement: The joint distribution is a "factored" distribution in which some random variables are either independent of or conditionally independent of most others.

Univ. of Wash.

The Joint Distribution

8

Jd
joint
Distribution

Why Bayes Networks?

Reasoning about events involving many parts or contingencies generally requires that a joint probability distribution be known. Such a distribution might require thousands of parameters. Modeling at this level of detail is typically not practical.

Bayes Nets require making assumptions about the relevance of some conditions to others. Once the assumptions are made, the joint distribution can be "factored" so that there are many fewer separate parameters that must be specified.


Univ. of Wash.

The Joint Distribution

9

Jd
joint
Distribution

Bayes Net for the Radio Problem



S: [T, F] – "A solar storm is happening."

B: [T, F] – "Mike's battery is on the fritz."

J: [T, F] – "Jan's radio works just fine."

R: [bad, good, none] – "Mike's radio reception"

J is independent of B

J and R are conditionally independent, both conditioned on S.

Univ. of Wash.


The Joint Distribution

10

Conditionally independent: given that C happens, events A,B are conditionally independent given C iff knowing that A happened, given C gives us no information on whether or not B happened given C. Mathematically:
 $P(A \text{ and } B | C) = P(A | C) P(B | C)$

Jd
joint
Distribution

Factored Distribution



S: $P(T), P(F) = [0.2, 0.8]$

B: $P(T), P(F) = [0.025, 0.975]$

J: $P(T|S=T), P(F|S=T), P(T|S=F), P(F|S=F) = [0.4, 0.6, 0.9, 0.1]$

R: $P(\text{bad}|S=T, B=T), P(\text{good}|S=T, B=T), P(\text{none}|S=T, B=T), P(\text{bad}|S=T, B=F), \dots = \dots$

This factored distribution uses 20 parameters, rather than 24 for the unfactored version.

Not all of these are independent parameters: By using $\sum p_i = 1$, we can reduce the numbers to 12 and 23. For larger numbers of nodes, the savings are often much greater.


Univ. of Wash.

The Joint Distribution

11

Jd
joint
Distribution

Working with the Bayes Net



S: **Solar Storm** (A solar storm is happening.)

J: **Jan's Radio** (Jan's radio works just fine.)

R: **Reception** (Mike's radio's reception).


Univ. of Wash.

The Joint Distribution

12

Jd
Joint
Distribution

Forward Propagation (from causes to effects)



Suppose S: there is a solar storm.

Then $P(J|S)$ is 0.4, $P(R=\text{bad}|S) = 0.493$, etc.

Suppose $\sim S$: no solar storm.

Then $P(J|\sim S)$ is 0.9, $P(R=\text{bad}|\sim S) = 0.203$, etc.

(These come directly from the given information.)


Univ. of Wash.

The Joint Distribution

13

Jd
Joint
Distribution

Marginal Probabilities (using forward propagation)



Then $P(J)$, the probability that J is true in any situation, is

$$P(J) = P(J|S)P(S) + P(J|\sim S)(1-P(S)) = 0.08 + 0.72 = 0.8$$

And $P(R=\text{bad})$, the prob. that R is bad in any situation, is

$$P(R=\text{bad}) = P(R=\text{bad}|S)P(S) + P(R=\text{bad}|\sim S)(1-P(S)) = (0.493)(0.2) + (0.203)(0.8) = 0.261$$

Marginalizing means eliminating a contingency by summing the probabilities for its different cases (here S and $\sim S$).


Univ. of Wash.

The Joint Distribution

14

Jd
Joint
Distribution

Using Bayes' Rule



Suppose we know J: Jan's radio works just fine.

How do we update the probability of S?

Bayes' rule: $P(S|J) = P(J|S)P(S)/P(J) = 0.08/0.8 = 0.1$


Univ. of Wash.

The Joint Distribution

15

Jd
Joint
Distribution

Updating Probabilities of Consequences



Suppose we know Jan's radio works just fine.

How do we update the probability of R=bad?

Use the revised probability of S:

$$P(R=\text{bad}|J) = P(R=\text{bad}|S)P(S|J) + P(R=\text{bad}|\sim S)P(\sim S|J) = (0.493)(0.1) + (0.203)(0.9) = 0.049 + 0.183 = 0.232$$

Which is slightly lower than $P(R=\text{bad}) = 0.261$.


Univ. of Wash.

The Joint Distribution

16

Jd
Joint
Distribution

Handling Multiple Causes



B: [T, F] – "Mike's battery is on the fritz."

$P(B) = 0.025$

| R | S | B | $P(R S,B)$ |
|--------|---|---|------------|
| R=bad | T | T | 0.75 |
| R=good | T | T | 0.2 |
| R=none | T | T | 0.05 |
| R=bad | T | F | 0.4 |
| R=good | T | F | 0.5 |
| R=none | T | F | 0.1 |
| R=bad | F | T | 0.6 |
| R=good | F | T | 0.3 |
| R=none | F | T | 0.1 |
| R=bad | F | F | 0.2 |
| R=good | F | F | 0.2 |
| R=none | F | F | 0.6 |


Univ. of Wash.

The Joint Distribution

17

Jd
Joint
Distribution

Explaining Away



Suppose R=bad. This raises the probability for each cause:
 $P(S|R=\text{bad}) = 0.378$, $P(B|R=\text{bad}) = P(R=\text{bad}|B)P(B)/P(R=\text{bad}) = 0.027$

Now, in addition, suppose not J (Jan's radio not fine).

Not J makes it more likely that S is true,

"And this explains R=bad."

B is now less probable: $P(B|R=\text{bad}, J=\text{F}) = 0.016$.

| R | S | B | $P(R S,B)$ |
|--------|---|---|------------|
| R=bad | T | T | 0.75 |
| R=good | T | T | 0.2 |
| R=none | T | T | 0.05 |
| R=bad | T | F | 0.4 |
| R=good | T | F | 0.5 |
| R=none | T | F | 0.1 |
| R=bad | F | T | 0.6 |
| R=good | F | T | 0.3 |
| R=none | F | T | 0.1 |
| R=bad | F | F | 0.2 |
| R=good | F | F | 0.2 |
| R=none | F | F | 0.6 |

Univ. of Wash.

The Joint Distribution

18



Benefits of Bayes Nets

“The decomposition of large probabilistic domains into weakly connected subsets through conditional independence is one of the most important developments in the recent history of AI.”

(Russell & Norvig, 3e. p499.)

Univ. of Wash.

The Joint Distribution

19



Benefits of Bayes Nets

The joint probability distribution with boolean random variables normally requires $2^n - 1$ independent parameters.

With Bayes Nets we only specify these parameters:

1. “root” node probabilities.
e. g., $P(A=\text{true}) = 0.2$; $P(A=\text{false})=0.8$.
2. For each non-root node, a table of 2^k values, where k is the number of parents of that node.
Typically $k < 5$.
3. Propagating probabilities happens along the paths in the net. With a full joint prob. dist., many more computations may be needed.

Univ. of Wash.

The Joint Distribution

20